

Seedance 2.5

The parts the launch pages left out.

ByteDance shipped a 30-second one-take video model with 50 reference slots and native audio. It is not a straight upgrade over 2.0 — it trades resolution for length and control. Here is the prompt grammar, a price list measured against the live API, and the one prompt clause that will silently destroy your renders after you have paid for them.

All firsthand testing was run on **OpenArt's** Seedance 2.5 endpoint — 13 generations, including a five-arm controlled experiment. Figures and available options are stated for that platform; other surfaces may differ.

MEASURED

I ran it against **OpenArt's live Seedance 2.5 endpoint** and read the result back.

REPORTED

Consistent across published sources. Not independently confirmed.

Every measured figure comes from one platform — OpenArt — on 8 August 2026. Other surfaces (Dreamina, Volcano Ark, fal, Higgsfield) may expose different options and price differently. Where that matters, it is flagged in the text.

00 Cheat sheet — one screen

01 The resolution nobody mentions

02 What it actually costs

03 The six-part formula

04 The four brackets

05 The accent trap

06 Reference tags

07 The depth map trick

08 The mistake everyone makes

The whole manual on one screen

SEEDANCE 2.5 — OPERATING CARD

Resolution	480p / 720p on OpenArt. No 4K there. Default is 480p — change it.
Frame rate	24 fps
Duration	4–30s native; 180s beta long mode
Prompt ceiling	30,000 characters
References	30 images + 10 videos + 10 audio = 50
Audio	Free. Never turn it off.
Music	Never let it generate music. Trips a copyright block after you pay.
Pricing	Linear per second. 30s = six 5s clips exactly. (OpenArt credits.)
Formula	Subject → Action → Scene → Style → Camera → Audio
Brackets	() music · < > SFX · { } dialogue · 【 】 subtitles
Camera	shot size → angle → movement. “Cinematic” is not a move.
Motion transfer	Use a depth map , not source footage.
Identity	Use a character sheet , not a hero shot.
Biggest mistake	Trying to oneshot the whole video.
Loop	Prototype 480p → lock seed → finish 720p → upscale

The “4K & 30s” headline does not survive contact with the API

I pulled the schema from **OpenArt’s** live Seedance 2.5 endpoint. The `resolution` field is an enum, and it contains exactly two values. **MEASURED**

```
"resolution": { "enum": ["480p", "720p"], "default": "480p" }
```

No 1080p. No 4K. And the default is **480p** — so everyone who never touched that dropdown has been generating at the lowest setting available.

It gets sharper. I checked an older **Seedance 2.0** job on the same account. It returned **3840×2160** — genuine 4K, from the older model, on the same platform.

MEASURED

SCOPE THIS CLAIM BEFORE YOU REPEAT IT

What I verified is that **OpenArt exposes 480p and 720p only**. I did not test Dreamina, Volcano Ark, fal or Higgsfield, so I cannot tell you 4K exists nowhere — only that it did not exist where I looked, while the same account got 4K out of Seedance 2.0.

That distinction matters. Say “no 4K anywhere” and one reply showing 1080p on another surface takes down everything else you claimed. Say “no 4K on the endpoint I tested, and 2.0 gave me 4K on that same endpoint” and the finding holds. Independent reports of a 720p ceiling and of Runway pointing users back to 2.0 for 1080p line up with this, but they are **REPORTED**, not mine.

Seedance 2.5 outputs at lower resolution than Seedance 2.0, and costs 62.5% more per second to do it.

That isn’t a scandal, it’s a trade: you are buying length and control, not pixels. But you should know you’re making it. Runway’s own docs reportedly tell users to switch back to 2.0 when they need 1080p — which tells you ByteDance knows exactly what this model is for. **REPORTED**

WORKING RULE

- Generate at 720p, always. Change the default the moment you open the tool.
- Budget a separate upscale pass. For client work it is not optional.

02 ECONOMICS

Measured prices, not estimates

Every pricing figure in circulation is an estimate. These were priced against

OpenArt's live endpoint on 8 August 2026, in **OpenArt credits**. **MEASURED**

Credits are a platform currency, so the absolute numbers only apply to OpenArt. The *ratios* are the portable part — and they are what the conclusions rest on. As a sanity check, Volcano Ark's published examples show roughly the same ~2.2× step from 480p to 720p. **REPORTED**

PRICE LIST — OPENART CREDITS · SEEDANCE 2.5 UNLESS NOTED

CONFIGURATION	CREDITS	PER SECOND
5s / 480p	300	60.0
30s / 480p	1,810	60.3
5s / 720p	650	130.0
30s / 720p	3,905	130.2
5s / 720p — audio OFF	650	130.0
Seedance 2.0 — 5s / 720p	400	80.0

Three things fall out of that table, and each one changes how you work.

Pricing is perfectly linear

A 30-second generation costs the same as six 5-second ones. No volume discount, no long-clip penalty. So choose duration for *story* reasons only — cost is not an argument in this decision.

Audio is free

720p/5s costs 650 credits whether audio is on or off. If you have been switching audio off to save money, you saved nothing and threw away the model’s best feature. Leave it on. Always.

720p is 2.17× the price of 480p

Worth it — and it makes the working loop obvious: prototype at 480p, finish at 720p.

03 GRAMMAR

ByteDance published an actual formula

Six slots, in this order. **REPORTED**

Subject → Action / Event → Scene & Environment →
Visual Style → Camera Movement or Cut → Audio

Most people write slot 1 and slot 3, then wonder why the camera does nothing.

The order matters because the model reads it as a priority stack. Put what you care about most earliest — and never contradict an earlier slot with a later one. Contradiction, not length, is what breaks these prompts.

04 AUDIO SYNTAX

Four brackets, four output channels

The model is said to route bracket types to different channels. **REPORTED**

BRACKET	ROUTES TO
()	Music
< >	Sound effects
{ }	Dialogue
[]	Subtitles / on-screen text

So instead of a mushy sentence about “cinematic ambient audio”, you write:

```
<heavy swell slapping the steel hull, wind buffeting the mic, distant gulls>  
(sparse low cello drone, single sustained note, no melody)  
{The sea does not care what you planned today.}
```

Three channels, separately addressed, no ambiguity about which is which.

HONESTY NOTE — I TESTED THIS AND IT DID NOT HOLD UP

Prompts using these brackets are **accepted without error**. But I ran a probe designed to detect whether `( )` actually routes to a dedicated music channel: I requested identical music in plain prose with no parentheses, holding everything else byte-identical.

If the bracket routed music somewhere prose could not reach, that run should have behaved differently. **It behaved identically**. The model responded to the *intent* to have music, not to the syntax.

That is one probe against one bracket, so it does not disprove the system — but it is the only firsthand evidence I have, and it points away from the claim. Treat the four-bracket table as a **useful way to organise your prompt**, which it genuinely is, rather than as a documented routing mechanism.

05 VOICE

Name the accent, or the model invents one

Native audio is the reason to use this model. Two rules. **REPORTED**

Name the regional variety, not the language

“English” gets you a neutral read. “Authentic Los Angeles English” gets you a performance. There is a widely-repeated case where a character kept receiving an unrequested British accent: `"American"` fixed it, and `"American English"` did **not**. The model is more sensitive to how you name the variety than seems reasonable.

Voice is inferred from your image references

Supply a character reference and say nothing about voice, and the model guesses what that face sounds like. It is frequently wrong and always confident.

Dialogue language: [language + regional variety].
The [character] says, [delivery direction]: {the actual line}

Keep lines **short**. Lip-sync in 2.5 is good enough that a short line lands cleanly and a long one exposes it.

06 REFERENCES

Fifty slots is a ceiling, not a target

Seedance 2.5 takes **50 references**: 30 images, 10 videos, 10 audio clips. Each file gets a handle — `@Image1`, `@Video1`, `@Audio1`. **REPORTED**

Rename the handles to roles

`@Image1` tells the model nothing. `@lead_character`, `@urban_night_style`, `@slow_dolly_motion` tell it a job. On API surfaces this is the `label` field on each reference object — I confirmed it exists in the live schema. **MEASURED**

Give every reference exactly one job, in a sentence

`@lead_character` provides Nia's face, hairstyle and jacket only.
`@urban_night_style` provides the colour palette and lighting, not the subject.
`@slow_dolly_motion` provides the camera path and pacing, not the content.
`@dramatic_voice` provides Nia's speaking voice.

Never let two assets define the same property. Two references both claiming “the style” is the most reliable way to get mush.

DO NOT FILL THE BUDGET

- Past a point, extra references **dilute** strong features and add noise.
- Conflicting assets weaken direction rather than strengthening it.
- Video and audio references work best at **5-10 seconds**, even though 30s is allowed.

Multi-keyframe storyboarding

You can hand the model an ordered storyboard:

Use @Image 1 through @Image 5 as keyframes in this order.
@Image 1 — she stands at the door, hand on the frame.
@Image 2 — she is halfway down the corridor, lights failing behind her.

Upload the frames **individually** — separate images align better than a collaged contact sheet. Keyframes control stage *order*, not literal frame reproduction: they are waypoints, not a flipbook.

07 DEPTH MAPS

The best-kept secret in the model

This is the most useful technique I found, and almost nobody outside a handful of creators is using it. **REPORTED**

The problem it solves: you want to transfer a motion — a dance, a walk, a specific camera move — onto your own character. So you hand Seedance the source video as a reference. And the choreography drifts.

Feed the model a video of someone dancing and it tries to interpret everything — the person, the clothes, the lighting, the room, and somewhere in there, the movement. Feed it a depth map and there's nothing left to interpret but the motion.

YOU HANDED IT TEN VARIABLES WHEN YOU MEANT TO HAND IT ONE

A depth map is motion and geometry with **every other variable stripped out**. No identity, no wardrobe, no lighting, no colour. It is the purest possible expression of “this reference has exactly one job” — the entire design principle of the reference system.

The workflow

- Build your storyboard normally, in GPT Image 2 or Nano Banana 2.
- **Convert that storyboard into a depth map** using the same image model.

- Hand Seedance three references, each with exactly one job:
 - the **depth map** → composition, framing, camera geometry, motion
 - a **style reference** → tone, palette, lighting
 - a **character sheet** → identity

Reported result: Seedance “understands AI-generated depth maps surprisingly well, giving you much more control with fewer tokens and less need to describe camera angles and compositions.”

That last clause is the kicker. **A depth map replaces paragraphs of camera description.** You stop writing “slow orbit from a low angle rising to eye level” and just show it the geometry. Fewer words, less contradiction, less drift.

The character sheet pipeline

Multiple creators converged on this independently. Nobody serious generates a hero shot and hopes it holds across a sequence.

- Create the character once.
- Turn it into a **reference sheet**: multiple angles, close-up portraits, clothing, accessories, textures, colours, proportions — instructing the image model to keep design and visual direction identical to the source.
- Feed the *sheet*, not the hero shot, as your identity reference.

A sheet gives the model a structured, readable breakdown and forces it to focus on essential visual information rather than stylistic noise. For a multi-scene sequence this is the difference between a character holding and their face quietly becoming someone else by shot four.

08 METHOD

The mistake almost everyone is making

Almost everyone prompting Seedance 2.5 right now is making the same mistake: trying to oneshot the whole video.

RECURRING OBSERVATION FROM CREATORS SHIPPING WITH 2.5

Thirty-second takes with fifty locked references **push the work upstream, into preparation.** Building the reference set, writing the brief, deciding the shot before you generate it. That is pre-production — a healthier place for the work to live than the retry button.

The principle that follows: **identity, wardrobe, props, location and camera should each have a clear source**, instead of fighting each other inside one

paragraph of prose. If all five are competing in your prompt text, you are running a 2026 model on a 2024 workflow.

Write beats, not paragraphs

Divide the video into consecutive stages. Give each stage **one main change and a visible end state**.

Timing: 0-5s hook – establish subject and world
 5-16s development – the thing starts happening
 16-25s escalation or proof
 25-30s resolution – land it

Seedance 2.0 rewarded dense, cut-heavy prompts. **2.5 punishes them**. 2.0-tuned prompts reportedly produce glitchy output on 2.5, and the model resists rapid one-second-cut styles. Scenes told to breathe do better. REPORTED

Write ten actions into a four-second clip and you get chaotic morphing. That is not a model failure. That is an instruction failure.

09 CAMERA

“Cinematic” is not a camera move

Specify three things, always, in this order. REPORTED

SLOT	OPTIONS
Shot size	wide, medium, close-up, macro, extreme close-up
Angle	eye-level, low angle, high angle, overhead, dutch
Movement	locked-off, dolly-in, push-in, pull-back, pan left, tracking, orbit, crane up, handheld, FPV

Terms reported to work reliably: *slow dolly-in, orbit, handheld, crane up, locked-off, smooth dolly forward, slow pan left, low tracking shot, pull-back, slow push-in, rack focus*.

The rule underneath: a slow push-in *feels* cinematic; a locked-off static shot feels like a product demo. The word “cinematic” gets you neither — it is a vibe word the model cannot act on. Every vague style adjective (“stunning”, “beautiful”, “photorealistic”) is dead weight. Replace them with things a camera operator could

actually execute: **rim lighting, shallow depth of field, slow push-in, low tracking shot.**

And give the move a **reason**. A push-in that lands on a reaction is direction. A push-in because push-ins look nice is drift.

10 MODES

Text-to-video is the least interesting thing it does

- **Keyframe** — supply a first frame *and* a last frame, prompt the motion between. This is how you make a shot end exactly where you need it.
- **Extend** — continue a clip **forwards or backwards**. Backward extension is the underused one: generate the moment *before* a shot you already like. It also bridges two existing clips.
- **Smart Edit** — change one element in place, region-level, without regenerating the shot. Point at the object or the face, fix that, leave the rest untouched.
- **Long-video mode (beta)** — up to **180 seconds**, currently via Dreamina.

WHY YOUR SETTINGS ARE GREYED OUT

- **Edit mode** locks aspect ratio to the input, and duration to input length $\pm 0.3s$.
- **First-frame mode** locks aspect ratio to your image, but leaves duration free.

11 MODERATION

The trap that eats your credits

I hit this one personally, and it is the most expensive gotcha in the model.

I submitted an entirely benign prompt — an Icelandic fisherman on a trawler delivering one line of dialogue. No violence, no public figure, no brand, no IP reference. It generated, and *then* failed. **MEASURED**

```
errorDisplayCode: "1003: output_moderation_blocked"
```

```
[byte-plus] output content rejected by moderation:  
OutputVideoSensitiveContentDetected.PolicyViolation:  
The request failed because the output video may be  
related to copyright restrictions.
```

Seedance runs **three separate moderation stages**: **REPORTED**

- a **prompt filter** on your text, before generation
- an **upload filter** — computer vision over your reference images, hunting copyrighted visual elements
- an **output filter** that reviews the finished video

Stage three is the one that hurts. Moderation runs during and after inference, so the GPU time is spent and **the credits are charged before the rejection fires**. Several platforms do not refund it. You pay full price for nothing.

The culprit is music — and it is provable

I ran a five-arm controlled experiment. Same seed, same duration, same resolution, one variable at a time. **MEASURED**

ARM	MUSIC	DIALOGUE	RESULT
Original	in ()	yes	BLOCKED
C1	in ()	yes	BLOCKED
C2	in ()	no	BLOCKED
C3	none	yes	RENDERED
D1	plain prose	yes	BLOCKED

- **The block is deterministic.** C1 was an identical resubmission and failed identically. A prompt that trips this trips it every time — you cannot retry your way past it.

- **Dialogue is innocent.** C2 removed the spoken line entirely and still failed.
- **Prompt-requested music is the trigger.** C3 kept the dialogue, the face, the seed and the sound effects, and removed only the music clause. It rendered clean.
- **It is not a syntax problem.** D1 asked for the same music in plain prose, no parentheses, and failed anyway. The model reads the *intent* to have music, not the bracket. You cannot rephrase your way out.

The mechanism: the model generates a score, an audio fingerprint matcher compares it against copyrighted material, and a match destroys the finished video *after* you have paid for it.

Context: in February 2026 Hollywood studios formally accused ByteDance of enabling mass copyright infringement via Seedance 2.0. ByteDance publicly committed to tightening controls. Things that passed early in 2026 are blocked now.

TWO RULES THAT FOLLOW

Never let the model generate music for work you care about. Write `no music of any kind, diegetic sound only` and lay music in during the edit, where you control the licence. Cleaner mix, and you stop tripping the matcher.

Smoke-test borderline concepts at 4s/480p (240 credits) before committing to 30s/720p (3,905 credits). A 16× cheaper failure is still a failure, but a survivable one.

Hard specs, read from the live schema

PARAMETER	VALUE	GRADE
Prompt length	up to 30,000 characters	MEASURED
Duration	4–30s integer; beta long mode to 180s	MEASURED
Resolution	480p , 720p — default 480p	MEASURED
480p resolves to	854 × 480	MEASURED
720p 16:9 resolves to	1280 × 720	MEASURED
720p 9:16 resolves to	720 × 1280	MEASURED
720p 21:9 resolves to	1470 × 630	MEASURED
Frame rate	24 fps	MEASURED
Aspect ratios	16:9, 4:3, 1:1, 3:4, 9:16, 21:9, adaptive	MEASURED
Audio	on by default, no extra cost	MEASURED
Seed	-1 random; set it to reproduce	MEASURED
Outputs per job	1–4	MEASURED
References	30 images + 10 videos + 10 audio = 50	REPORTED

Note the prompt ceiling: **30,000 characters**. Every guide telling you to keep prompts short is giving you a style opinion, not a technical limit.

And note **24 fps**. If you are cutting this into a 30 or 60fps timeline, conform it deliberately instead of letting your editor guess.

Known weaknesses — plan around these

- **Physical plausibility in complex motion.** Collisions, cloth, liquids under force — that is where it breaks.
- **Many interacting subjects.** Stability degrades as subject count rises. Two people talking is fine; a crowd is a gamble.

- **Fast action morphs.** Decoherence in rapid sequences persists.
- **Generation stalls.** Normal is 2–3 minutes. I saw 8s/720p jobs exceed **20 minutes**, and one pass 30, on the day the API went wide. **MEASURED**

The shot this model wins at: **one clear subject, doing one clear thing, with a camera move that has a reason.**

13 FIELD

When to reach for something else

There is no universal winner, and anyone telling you there is has something to sell.
Pick per shot. **REPORTED**

MODEL	WINS AT	REACH FOR IT WHEN
Seedance 2.5	Duration, reference control. The <i>only</i> model accepting audio reference input .	Character consistency across shots, long takes, ads, anything reference-heavy.
Veo 3.1	Cinematic grade, film-like motion blur, professional lighting, strong native audio.	The hero shot, where picture quality is non-negotiable.
Kling 3.0	Raw resolution — 4K/60fps — plus a free tier.	Volume social output and prototyping on a budget.
Sora 2	Physics, spatial awareness, eye contact and hand gestures.	Nothing new — see below.

DO NOT START A PIPELINE ON SORA 2

OpenAI deprecated the Videos API and the Sora 2 models on **24 March 2026** and published a shutdown date of **24 September 2026** — 47 days from this revision.

Sora 2 still wins on physics, and it will keep topping comparison posts written by people who never checked the deprecation notice. Do not build anything on a model with a published end-of-life inside two months.

The practical split: **Seedance 2.5** when a character has to survive more than one shot, **Veo 3.1** for the one frame that has to look expensive, **Kling 3.0** for volume.

Copy-paste starting points

Cinematic dialogue shot

[Character: concrete identity anchor – age, build, one memorable garment, one facial detail]. [What they do]. [Where, with time of day and weather].

Camera: [one move with a reason – push-in, orbit, handheld follow], [lens], [depth of field].

Timing: 0-2s [beat]. 2-4s [beat]. 4-6s [beat].

Dialogue language: [language + regional variety], [delivery].
The [character] says: {short line}

<diegetic sound 1, diegetic sound 2, room tone>

Constraints: one continuous shot, no cuts, no on-screen text, no music of any kind, diegetic sound only.

Product / macro

[Product] on [surface]. [The one thing that moves].

Camera: [locked-off, then one move], [macro lens], very shallow depth of field, rack focus from [A] to [B].

Timing: 0-2s stillness. 2-4s [entry]. 4-6s [the move].

<one or two precise material sounds>

Constraints: single take, no cuts, no text, no logo, [lighting], no lens flare, no music.

30-second narrative beat sheet

[Subject + world in one sentence.]

Beat 1 (0-5s): hook. [What we see first.]

Beat 2 (5-16s): development. [What starts happening.]

Beat 3 (16-25s): escalation. [The turn.]

Beat 4 (25-30s): resolution. [Where it lands.]

Camera through-line: [one continuous logic across all four beats].

@lead_character provides [exactly one property].

@style_ref provides [exactly one property].

<ambience>

Constraints: one continuous take, let each beat breathe,
no rapid cuts, no music.

15 LOOP

The working loop

- **Prototype at 480p.** 60 credits/second. Iterate the prompt here, not at 720p.
- **Lock the seed** once a take is nearly right, then change one thing at a time.
- **Promote to 720p** only when composition and motion are settled.
- **Upscale as a separate pass.** Build it into the budget from day one.
- **Use Extend and Smart Edit** instead of regenerating. A take you like is expensive to rediscover.
- **Keep audio on** throughout — it is free, and it is the differentiator.
- **Keep music out** of the generation. Add it in the edit.

How this was researched. The schema, the resolution enum and 480p default, the 24fps output, the frame dimensions, the full price list, the audio-is-free finding, the 2.0-outputs-4K comparison and the five-arm moderation experiment were all verified firsthand on **OpenArt's** Seedance 2.5 endpoint ([byte-plus-seedance-2-5](#)) on 8 August 2026 — 13 generations in total.

What that scopes. OpenArt serves ByteDance's model, it is not ByteDance. Available resolutions, credit prices and moderation behaviour are what OpenArt exposed on that date. The moderation block originates upstream — the error string is returned verbatim from `byte-plus` — so that finding should travel across platforms. The resolution ceiling and the credit prices are the ones most likely to differ elsewhere, and are flagged where they appear.

Everything marked `REPORTED` — the bracket syntax, the accent behaviour, the `@tag` semantics, the modes and parameter locks — comes from published guides and creator reports that could not be opened directly, because the research environment blocked outbound page fetches. It is consistent across multiple independent sources, which is why it is here. It is not the same grade of evidence as the measured numbers.

A guide with its seams showing beats one that sounds certain and isn't.